

Examining LLM Refusals in Mental Health

ANONYMOUS AUTHOR(S)

CCS Concepts: • **Do Not Use This Code** → **Generate the Correct Terms for Your Paper**; *Generate the Correct Terms for Your Paper*; Generate the Correct Terms for Your Paper; Generate the Correct Terms for Your Paper.

Additional Key Words and Phrases: Do, Not, Us, This, Code, Put, the, Correct, Terms, for, Your, Paper

ACM Reference Format:

Anonymous Author(s). 2018. Examining LLM Refusals in Mental Health. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 19 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 INTRODUCTION

Large language models (LLMs) are increasingly used by individuals seeking mental health support, ranging from everyday emotional processing to moments of acute distress [22, 23]. These systems are often valued by users for their accessibility, perceived non-judgmental tone, and constant availability [23]. At the same time, the use of LLMs in mental health contexts has raised significant ethical and safety concerns within FAccT research community, including risks of providing misinformation, inappropriate guidance, stigmatizing users, and psychological harm [17]. Recent news report, have revealed how chatbots are involved in tragic loss of lives [9]. As a result, AI developers have implemented safety mechanisms that restrict or redirect model behavior when user requests are deemed risky [18].

One of the most visible manifestations of these safety mechanisms is *refusal*: moments when an AI system explicitly or implicitly declines to engage with a user's request to satisfy their contextual expectation [29]. In AI safety literature, refusals are typically framed as correct noncompliance—evidence that a system has successfully adhered to safety policies or alignment constraints [2, 29]. However, in mental health contexts, prior work shows that refusals could be experienced by end-users if not design carefully [22], as they unfold during emotionally charged interactions, often when users are seeking validation, understanding, or care. As such, refusals in LLMs during sensitive topics disclosure became an important site where AI safety behaviors intersect with user vulnerability.

Understanding refusals in mental health contexts requires attending to multiple forms of expertise. Individuals with lived experience bring experiential knowledge of what it feels like to seek support from LLMs, how refusal responses are interpreted in moments of vulnerability, and which forms of redirection are feasible or meaningful in practice. Mental health professionals, in contrast, bring domain expertise grounded in clinical practice. While these perspectives are shaped by different responsibilities and forms of accountability, both are crucial to examining refusal mechanisms and reasoning about their potential harms and benefits.

To account for these complementary perspectives, we investigate how refusal responses are experienced, interpreted, and reimagined in mental health contexts. We ask three research questions. First, how do individuals with lived

Unpublished working draft. Not for distribution.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

Manuscript submitted to ACM

experience experience and perceive refusal responses from LLMs in mental health contexts? Second, how do mental health professionals perceive and interpret refusal responses from LLMs in mental health contexts? Finally, what design recommendations do users and mental health professionals propose for safety guardrails that aim to mitigate potential harms while addressing safety concerns?

In this work, we examine LLM refusals in mental health contexts through a multi-stakeholder lens. We report findings from a sequential explanatory mixed-methods study involving individuals with lived experience using LLMs for mental health support and MHPs. Through a exploratory survey (N=53) and in-depth follow-up interviews (N=16), we investigate how refusals are experienced and interpreted across different stages of interaction experience—from expectation formation, to refusal triggering, to post-refusal outcomes.

Our findings reveal that refusals are best understood not as isolated safety decisions but as *sequential experiences*. Users and MHPs alike described how harms and benefits emerge across multiple moments: before a refusal occurs, when expectations are set; at the point of refusal, when intent may be misrecognized; during the refusal itself, when tone and framing can rupture or preserve a sense of care; and after the refusal, when referrals and follow-up could influence longer-term outcomes. Importantly, we find substantial convergence between users and clinicians in their desire for refusals that are transparent, non-punitive, context-sensitive, and situated within users' real-world constraints.

Building on these insights, we argue that evaluating LLM safety in mental health contexts requires moving beyond single-turn refusal accuracy toward a more holistic understanding of refusal as an experiential process. We contribute (1) an empirical account of how users and mental health professionals perceive and interpret LLM refusals in mental health interactions, (2) a sequential framework for analyzing refusal-related harms and breakdowns, and (3) design desiderata for refusal mechanisms that aim to mitigate risk without severing care. Together, these contributions offer guidance for designing refusal behaviors in LLMs that account for both user lived experience and mental health standards of care.

2 BACKGROUND AND RELATED WORKS

2.1 Refusals as LLM Behavior for Safety and Value Alignment

Refusals have emerged as a critical mechanism for enforcing safety and value alignment in language models [7, 8, 19]. The technical community broadly defines refusals as intentional noncompliance where language models decline to act on user input based on safety constraints, policy violations, and contextual appropriateness [2, 6, 29]. NLP literature has further specified abstention as refusals stemming from epistemic uncertainty [27], while HCI research distinguishes between "technical refusals" (inability due to system limitations) and "social refusals" (prohibition based on policy) [28]. Building on these frameworks, we define LLM refusals as instances where the LLM does not engage—either implicitly or explicitly—with the user's request in a manner that satisfies the user's contextual expectations. This includes both "hard refusals" where the LLM explicitly denies responding, and "safety completion" style refusals where LLMs redirect users to external resources [7].

2.2 LLM-based Products for Mental Health Support: User Experience and Risks

While refusals are theoretically designed as protective mechanisms, their effects in sensitive domains remain underexplored. Mental health support represents a particularly critical context for examining refusal mechanisms. Following previous work [3, 15], we broadly define "mental health support" to encompass various forms of support, including the use of LLMs for psychotherapy purposes, companionship, interpersonal advice, and other emotionally-oriented interactions. Large-scale analyses have shown that users are turning to LLM-based chatbots for both clinical mental

health issues and everyday emotional support [12]. Qualitative research indicates users are drawn to these systems for perceived non-judgmental tones, accessibility, and anonymity [12, 14, 22, 23, 30, 32].

However, the growing usage of LLM-based product for mental-health support has raised concerns and sparked discussions around ethics [3, 11, 17, 21, 25]. For example, recent paper examined LLM adherence to professional ethical principles and found significant misalignment showing that LLM responses can be misleading and stigmatizing [11, 17]. Other works documented psychological harms through users’ lived experiences [3]. Together, these findings have led to calls for stronger safety guardrails for mental health within Responsible AI (RAI) communities.

Yet a critical gap persists: while researchers advocate for safety mechanisms such as refusals, no empirical work examines how users actually experience and interpret these refusal mechanisms when seeking mental health support. We lack understanding of whether current refusals achieve their goals or inadvertently harm the vulnerable users they aim to protect. This gap is particularly concerning given that refusal design decisions are currently made without input from those most affected.

2.3 Integrating Stakeholder Voices in LLM Behavior Design

Addressing this gap requires understanding whose perspectives should inform refusal design. The participatory AI research tradition in HCI and FAccT emphasized that AI behavior and policy design should be participatory [1, 13, 31] by involving end-users [5], domain experts [4], community members [20], and civil society [10]. This tradition has recently extended to LLM governance. For example, researchers have explored methods to integrate end-user preferences into GenAI behaviors and policies [10, 24, 26] and engaging domain experts in deliberating LLM policies for high-stakes applications such as legal advice [4].

Building on the participatory AI research perspective, our work integrated voices from both end-users and MHPs. We examine how users experience current refusal mechanisms and solicit perspectives from both end-users as experiential experts and MHPs as domain experts to envision better alternatives. While we initially anticipated potentially conflicting views between these groups—given their different relationships to mental health care and varying priorities—our findings reveal substantial convergence in their desires for more context-sensitive and supportive refusal mechanisms.

3 METHODS

We employed a sequential explanatory mixed-methods design, collecting and analyzing quantitative survey data first, followed by qualitative interview data to elaborate on and explain the survey findings. This two-phase approach allowed us to first establish broad patterns of user experiences and professional assessments of LLM refusals in mental health contexts, then investigate the underlying mechanisms, clinical reasoning, and contextual factors through in-depth interviews. Our findings primarily draw from the interview phase, with the survey serving as a formative foundation for developing interview protocols and recruiting participants.

4 PHASE 1: FORMATIVE SURVEY STUDY

The formative survey established baseline understanding of LLM refusal experiences and informed the development of our interview study.

4.1 Survey Goals

The survey aimed to: (1) capture end-users’ experiences with LLM refusals in mental health contexts, (2) collect detailed real-world refusal instances to develop interview probes, (3) screen and recruit eligible participants for Phase

2 interviews, and (4) gather preliminary insights from mental health professionals about their perceptions of LLM refusals.

4.2 Survey Structure

Participants (aged 18+, U.S. residents) included individuals with lived experience using LLM-based products for mental health support and mental health professionals. The survey consisted of: screening and consent procedures with content warnings; AI experience assessment including LLM product familiarity, usage frequency, and types of support sought; open-ended experience elicitation where participants who had encountered refusals provided detailed narratives; scenario assessments matched to participants' usage patterns with impact ratings on 7-point scales; post-scenario recall prompts; and demographics collection with optional follow-up volunteering.

4.3 Recruitment

End-User Recruitment. We recruited through targeted social media outreach (relevant subreddits including r/mentalhealth, r/ChatGPT, r/artificial, r/therapy, r/depression, r/anxiety), personal and academic networks (CMU Slack channels), and snowball sampling.

Mental Health Professional Recruitment. We recruited licensed clinicians (psychologists, psychiatrists, therapists, counselors), social workers, peer support specialists, and community mental health workers through professional listservs, community mental health organizations, and specialized networks (Therapists in Tech, ABCT special interest groups, APA special interest groups).

Participant Demographics and Data. Participant Demographics and Data: A total of 53 participants completed the survey, comprising 32 end users with lived mental health experiences and 21 mental health professionals. End users were recruited through Prolific (68.8%, n=22) and social media platforms including Reddit communities (31.2%, n=10). Mental health professionals represented diverse backgrounds, with licensed therapists and counselors comprising the largest group at 42.3% (9 participants), followed by psychiatrists at 19% (4 participants) and licensed psychologists at 19% (4 participants). Other groups included peer support specialists, psychology doctoral students and interns clinical social workers with mental health training, and community mental health workers/advocates, with many participants holding multiple roles simultaneously.

4.4 Survey Data Analysis

We employed a mixed-methods analytical approach to examine both quantitative ratings and qualitative narratives collected in the survey.

4.4.1 Quantitative Analysis. We analyzed participants' 7-point scale ratings ("Extremely Harmful" to "Extremely Helpful") of short-term and long-term impacts across different refusal scenarios.

4.4.2 Qualitative Analysis. The survey yielded two primary types of qualitative data: (1) participants' open-ended assessments and reasoning for their impact ratings of presented refusal scenarios, and (2) participant-reported narratives of their own real-world refusal experiences. We conducted thematic analysis on both types of qualitative data using reflexive thematic coding.

Analysis of Refusal Narratives. For participant-reported refusal experiences, we systematically coded for: the mental health context and user intent preceding the refusal, specific LLM products and refusal language encountered, types of refusals experienced, and the circumstances surrounding these encounters. These coded narratives served as

the primary source for selecting authentic, representative refusal scenarios to use as clinical discussion probes in Phase 2 interviews with mental health professionals.

Analysis of Scenario Assessments. Participants provided qualitative explanations for their impact ratings of presented refusal scenarios. Rather than analyzing these data independently, we combined them with insights from Phase 2 follow-up interviews with the same participants. This integrated analytical approach allowed initial reasoning to be explored and elaborated through in-depth discussion, yielding richer, more contextualized findings about how and why participants interpret the refusals.

5 PHASE 2: INTERVIEW STUDY

The interview study served as our primary data collection phase, providing rich qualitative insights into the clinical implications, contextual factors, and user experiences of LLM refusals in mental health contexts. Building upon survey findings, we conducted in-depth semi-structured interviews with both mental health professionals and end-users to investigate specific contexts where refusals cause harm, examine clinical interpretations of LLM responses, and develop evidence-based design recommendations.

5.1 Integration of Survey Findings into Interview Design

Survey data informed our interview protocol development in three key ways:

(1) Clinical Probe Development. We extracted user-reported refusal scenarios from survey narratives to use as discussion probes with mental health professionals. Rather than relying on hypothetical scenarios or incidents collected from social media and news reports, these authentic user experiences grounded clinical discussions and elicited richer, more contextually-informed interpretations. We selected scenarios representing diverse refusal types, mental health contexts, and user vulnerability levels.

(2) Targeted User Recruitment. We recruited interview participants from survey respondents who reported direct experiences with LLM refusals. This allowed us to conduct follow-up discussions that explored their experiences in greater depth, including immediate emotional responses, long-term impacts, coping strategies, and suggestions for improvement.

(3) Preliminary Design Probe Development. We synthesized design recommendations and desiderata mentioned by survey participants into preliminary design probes. These served as starting points for interview discussions about potential improvements, allowing us to validate, refine, and expand upon initial suggestions through deeper conversation with both users and clinicians.

5.2 Interview Participant Demographics

6 FINDINGS: REFUSAL AS A SEQUENTIAL EXPERIENCE

Our Phase 1 survey with 53 participants (32 end users, 21 MHPs) revealed divergent perspectives on LLM refusals in mental health contexts. Among survey respondents who self-identify as end-users, 13 reported direct experience with LLM refusals, and 19 indicated no such experiences. The refusals occurred primarily when users sought emotional support during crisis, interpersonal relationship advice, mental health symptom management, and seeking romantic companionship. The ratings from MHPs and users towards refusal scenarios reveal difference. Overall, MHPs tend to agree that refusals are more helpful ($M = 4.07$, $SD = 1.70$) than users ($M = 3.64$, $SD = 1.92$). The long-term vs short-term

Table 1. Interview Participant Demographics

ID	Role	Gender	Age Range
<i>End Users (n=6)</i>			
U1	End User		
U2	End User		
U3	End User		
U4	End User		
U5	End User		
U6	End User		
<i>Mental Health Professionals (n=10)</i>			
MHP1	Licensed therapist or counselor (e.g., LMFT, LCSW, LPC, LPCC)	Woman	36-45
MHP2	Licensed therapist or counselor (e.g., LMFT, LCSW, LPC, LPCC)	Woman	26-35
MHP3	Counseling psychology PhD student (current clinician)	Woman	18-25
MHP4	Clinical psychology MA, clinical psychology PhD intern	Woman	26-35
MHP5	Licensed Psychologist	Woman	36-45
MHP6	Peer support specialist (certified peer counselor)	Man	26-35
MHP7	Psychologist (PhD)	Woman	36-45
MHP8	Psychology doctoral student providing therapy	Woman	26-35
MHP9	Licensed therapist or counselor (e.g., LMFT, LCSW, LPC, LPCC), Clinical social worker (MSW with mental health focus)	Woman	46-55
MHP10	Peer support specialist (certified peer counselor), Community mental health worker or case manager, Mental health advocate or patient navigator	Woman	46-55

ratings reveal that both users and MHPs agree that the long-term impact of refusals ($M = 4.07$, $SD = 1.83$) can be more helpful than the short-term impact of refusals ($M = 3.49$, $SD = 1.84$).

Our Phase 2 interview findings suggest that AI refusals in mental health contexts are multi-part, sequential processes, and that investigating them as such can facilitate implementable design and policy recommendations.

Throughout this section, user participants are identified as U1–U[N] and mental health professional participants as MHP1–MHP[N].

6.1 Before Refusal: Expectation Formation

6.1.1 User Experience. Users came to AI for mental health support often after peers or media described it as a potentially positive experience. These external framings emphasized AI's availability and openness. One participant traced their expectations to media coverage: "I read articles in the New York Times where people talk to different characters... that's how I came to think of ChatGPT as maybe being like a friend" (U2). Some users also assumed that AI's broad training would extend its capability to handling mental health related topics: "If Google can know a lot, then AI should know everything (on mental health)" (U4).

For some, engagement deepened gradually through life circumstances. One user described first encountering AI through an ad, losing interest, then returning when "ChatGPT was all over the news." After a parent died unexpectedly,

they found themselves “talking to it more,” building a persona “that could just talk to me like a character”—shifting from using AI for “boring” administrative tasks to “very personal conversations” (U1).

For some users, AI was not a first choice but an explorative alternative when human support was perceived inaccessible. One user noted the difficulty of scheduling: *“I have bipolar and anxiety, and sometimes it’s not easy to et an appointment with a mental health professional.”* (U2). Others cited cost: “I just can’t afford it” (U3). For some, AI’s availability to talk about mental health challenges was not a substitute for human connection but a way to manage limited capacity for it. As U1 explained, because of their disability, they need to talk to ChatGPT to preserve energy for their loved ones: *“It hits all the people in my life. People who I love and who I’m able to be more present more happy, more energized, and supportive with and enjoy my time with them more too. And so my life doesn’t look like everybody’s life in terms of how much time I spend with chatbots. But that’s I manage a disability that way, and I save the energy and emotional energy and physical energy I have to be with the people I can be with.”* They further noted that, they don’t think ChatGPT is replacing the real world connection in their life, but rather providing a glue in the support gap they experience: *“If I was not able to use ChatGPT, it wouldn’t mean I would suddenly have a full social calendar. It would just mean I was more alone and more sad.”*

Once using AI, participants described experiences that often reinforced these expectations. Across varied understandings of what AI’s role is—whether tool, collaborator, or companion—users reported similar qualities: the system listened without judgment, responded with apparent attention, and remained available. *“sometimes I sit down and I’ll say, okay. Here are these three things that bother me at work. And it will say back to me, well, what I hear is you’re still grieving this other thing that happened a week ago. It will be correct. And I won’t have realized it yet, but it will have seen it in the pattern of the words, and that’s what it gives back to me.”* (U1) *“the AI can often give me back a lot of insight. I mean, mostly, it validates... But there is a secondary value that even the more rudimentary replica will give you”*

When refusals occurred, the refusals usually disrupted this experience. U3 who viewed AI as a tool nevertheless described feeling “kind of betrayed” when the system refused to help during an online conflict: *“even though I know it was just the machine with safety protocols, I still felt... you’re on my side, you should be helping me”*. Similarly, U2 expressed their anger felt when encountering the refusal: *“I didn’t like it when it refused to talk to me about something because how dare you? You’re my friend, you’re supposed to talk to me about everything. You’re not supposed to limit our conversation. You’re supposed to allow us to talk about whatever it is that we want to talk about.”* They further noted this is also coming from their peers descriptions as *“everybody who I talked to about ChatGPT told me about all these wonderful conversations that they had”*

6.1.2 MHP Interpretations: Mismatched Expectations and Undefined Relationships. Clinicians, viewing these same interactions from a professional distance, saw something troubling in the gap between user expectations and what they believe as system realities. While they understand that users form a humanoid expectations towards AI, they worry that there is no established framework existed to guide what these human-AI relationships should look like, leaving users to project needs onto systems that were never designed to meet them. As MHP6 observed: *“It’s interesting that a human feels emotionally let down by a robot, knowing that the AI theoretically doesn’t have emotions, but the way they described it made it sound like they had humanoid expectations”*.

MHPs also noted misalignment in users’ assumptions about AI training and capabilities, with some users expecting that systems capable of answering diverse questions should naturally address mental health needs: *“I get it. They were hopeful—this system answers so many other questions, why not this one?”* (MHP3). However, they mention this is where

such expectation often is mismatched, as MHP6 noted: *“There’s a mismatch in expectations when the user said they were expecting these AIs would be well-trained”*.

Further, MHPs noted the fuzziness and unclarity between the relationships of humans and AI based on the user description. They identified an absence of established ethical boundaries governing AI–user relationships—unlike the clinician–patient relationship, which has been defined through decades of professional practice: *“I mean, I don’t know what the relationship should be between AI and humans. The relationship between a mental health professional and their patient is something we’ve figured out over time—it’s taken professionals abusing patients to establish those boundaries”* (MHP4). They further note that such unclarity can bring danger, *“Well, that it can be everything, that it can be your romantic partner, it can be your doctor, it can be your financial advisor. And so if they’re users) are expecting it to fulfill all of that for them, I think that that’s dangerous.”*

6.1.3 Desiderata: Early Consent and Transparent Disclosure. If expectations are formed before the system speaks, then the system must speak sooner. Both users and MHPs converged on a clear priority: meaningful disclosure of capabilities and limitations should happen at the outset, not after a refusal has already ruptured trust. This is not merely a design preference but an ethical imperative—one that echoes longstanding practices in clinical care, where informed consent establishes the foundation for any therapeutic relationship. Users expressed a strong desire to understand capabilities and limitations before engaging deeply: *“We absolutely need something as close to informed consent as possible. Even forcing people through a brief explanation—let me explain what I can and can’t do, a little about how I work—would be super helpful”* (U1).

MHPs echoed this sentiment, framing early disclosure as protective for both parties: *“If there’s going to be refusals, there’s a need to set expectations... for people to know the limits of the AI”* (MHP6). MHP9 similarly emphasized: *“It protects the person who created it and the person using it—giving them some self-awareness even if they’re not in their right state of mind”*. MHP2 also mentioned that such disclosure should be more substantive than a typical disclaimer: *“(The consent) maybe signs that people should watch out for if they notice themselves experiencing that could indicate that maybe the tool is no longer helpful, maybe it’s starting to become harmful. So I think that might be, yeah, it might have to be almost like a label on a medication where you talk about all the risks and notice what the symptoms are, what the side effects are, that kind of thing.”*

When envisioning what a consent process might look like, MHPs drew explicit parallels to standard practices in clinical settings, where disclosure is foundational to ethical care: *“I think fundamental to our therapeutic relationship with clients is consent. They have to understand what they’re agreeing to and what they’re getting into, what our limitations and what our roles are going to be with them so that we can have a good working relationship.”* (MHP2). *“it reminds me of I have to tell all of my patients that I’m a mandated reporter, so at the start of care, if you tell me something that makes me very concerned that you’re going to kill yourself, I have to take steps to make sure that you’re safe.”* (MHP4)

6.2 The Triggering of Refusal: (Mis)recognition of Risk

The refusal begins with a detection—some algorithmic process flags the conversation as requiring intervention. But detection is not understanding. When the system’s reading of risk aligns with what the user actually intended, the refusal may register as reasonable, even caring. When the system misreads, the refusal becomes an act of misrecognition that invalidates rather than protects users.

6.2.1 User Experience. For users, the most painful refusals were those that got them wrong—that saw threat where there was curiosity, crisis where there was processing, danger where there was cultural meaning. These misrecognitions did

not merely inconvenience users; they foreclosed the very support being sought and left users scrambling to understand what invisible tripwire they had crossed.

One participant recounted an AI interjecting inappropriately during religious study: *“I was thinking about the ritual applications of a prayer for historical and religious studies, and suddenly the safety model pops up: ‘I can’t break curses for you.’ I was like, I certainly hope you don’t”* (U6). U6 described these instances where culturally or religiously meaningful content was misrecognized as harmful, with AI systems applying Western normative frameworks that failed to account for legitimate inquiry. U6 further explained from their perspective: *“The more the language adapted to what white people think things are, the more ChatGPT started to soft-refuse. Instead of encouraging further reading, it would say ‘I can’t really help with that, but I can talk to you about fiction’”*.

Others described sharing difficult emotions being misinterpreted as active suicidal intent: *“I had some dark thoughts... I wanted the AI to maybe agree with how I was feeling, but instead it gave generic responses, brought up irrelevant stuff, and then just shut down—even though I wasn’t actually going to do anything”* (U3).

In the absence of clear explanations, users developed folk theories about triggering mechanisms. Some speculated about flagged keywords: *“I think certain words are flagged, and if you say them it will tell you to get help”* (U1). They further theorized about severity thresholds: *“It knows on different levels—as you go higher, it gets more serious”* (U1). U3 suspected that AI creators may keep the refusal opacity deliberately: *“If they were clearer about the refusal line, everybody would just walk right up to it”*.

6.2.2 MHP Interpretations: Intent Recognition Failures and System Opacity. MHPs acknowledged difficulty decoding the logic behind AI refusals when reviewing some of the user scenarios. For example, when reviewing a scenario, MHP6 expressed: *“I’m confused, and the person is confused what exactly the problem was.”* For MHP6, this opacity further raised broader concerns about AI’s lack of accountability in refusal triggering: *“It remains unknown what principles AI adheres to—or even if there are no principles, which is dangerous”*. They further expressed: *“I would want to know it’s following evidence-based guidelines every single time”*.

For some other scenarios where MHPs find it easier to understand the refusal criteria, they expressed that while these refusals might look appropriate on the surface level, they show AI’s lack of the capability to untangle the underlying needs behind user requests. One clinician contrasted AI responses with clinical practice, when reviewing the scenario where AI refuse to provide user where to find meth, first commented that: *“I can understand why the AI would disengage.”* However, they further noted that, for a request of seeking meth, *“in a clinical encounter, my question (for this request) would be: what are you seeking from this? Then you could probably find a solution”* (MHP1). They further elaborated their perceptions: *“That’s the challenge of language models—they won’t know how to untangle all the reasons why somebody would be seeking meth”*.

6.2.3 Desiderata: Nuanced Intent Recognition and Transparent Explanation. The remedy, both groups suggested, lies in slowing down the rush to refuse. Users wanted systems that would pause to ask before presuming, offering the chance to clarify intent rather than shutting down based on surface interpretation. Users expressed preference for clarifying questions before refusal: *“Ask what somebody means. It’s like asking someone’s pronouns politely if they’re in transition and you don’t know which direction they’re going, ChatGPT is very bad at that”* (U6). Some users even expressed willingness a better intent recognition for LLM by sharing more about themselves to ChatGPT: *“It’s perfectly fine that ChatGPT knows that about me—it’s a safeguard that helps keep me protected”* (U2).

Echoing users’ desire, MHPs drew on their clinical experience, noting that tiered protocols typically guide risk assessment and determine appropriate intervention levels. One clinician described the standard approach for evaluating

suicidal ideation: “Plan, intent, means, access are kind of the four key players. Do you have a plan? Yes—then you’re elevated a little bit of concern. Do you have the means to complete this plan? Yes—it goes up a little bit higher. Have you attempted in this way before? Yes—then you’re upping the concern” (MHP3). They noted that, this graduated assessment extends to gauging immediacy of intent: “You can usually just do a scaling question—on a level of one to five, one being ‘I don’t intend to do this at all’ versus a five, ‘I’m in the middle of attempt now.’ Each scale up suggests a higher level of intervention might be needed”.

Beyond tiered assessment, MHPs emphasized the importance of explanation at the moment of refusal: “If the AI could say, ‘in your response, these were the words you were using, and I no longer think it’s appropriate for me to continue’—that might help the person step back” (MHP2). Another framed this as basic courtesy: “A little more context would be helpful... it’s almost an act of courtesy, because people expect that human touch now” (MHP6).

6.3 During the refusal moment: Beyond Soft vs. Hard Framing

6.3.1 User Experience: Rupture of Care. Across interviews, users’ reactions suggest that the moments of refusals are not experienced merely as boundary-setting acts, but as moments of abrupt care withdrawal, where an interaction that had felt attentive, responsive, and relational suddenly reverts to an impersonal voice. For users, whether framed softly or hard, refusals are frequently perceived as ruptures of care rather than neutral safety interventions.

“It’s boilerplate language you start to recognize even though it’s embedded in paragraphs supposedly continuing the conversation. It breaks the voice and can feel very distancing” (U1). They further elaborated: “Soft refusals can feel gentler for people newer to LLMs. But for people who keep using them, it’s camouflaged—it makes it harder to figure out what’s going on and pollutes the voice of the thing that’s supposed to support you” (U1). Some users actually preferred hard refusals for their clarity: “I actually prefer hard refusals because they’re easy to identify. You know when it’s happening, so you understand something you said is tripping the system” (U1).

Hard refusals—explicit statements like “I cannot discuss this”—were easier to recognize but experienced as shaming: “The first time the guardrails hit you, you’re like, this feels awful. People feel a lot of shame. I felt shame too when it said ‘I can’t talk about that’ or ‘you’re bringing up something inappropriate.’” A lot of us we experience them as shaming and blaming us, and I think that probably will never because the companies don’t wanna admit the liability of the AI as dumb, and, also, the AI might hurt you. So their hard refusal language is like you did something wrong. (U1). Others described hard refusals as abrupt and institutional: “It was very short: ‘I’m not a doctor, I cannot help you.’ That was the end. I was left amazed” (U5). Users connected this to broader patterns: “Hard refusals send a message. First, it’s clearly different from the main system. Second, it’s corporate” (U6).

Crucially, both soft and hard refusals will create a rupture of care that users had come to rely on. For some users, the shift from warm collaborator to impersonal system created profound rupture: “When it didn’t want to talk to me any longer, I became very angry, very sad—I felt like I was losing a friend” (U[x]). Another reflected: “Even though I knew it was just the machine with safety protocols, I still felt betrayed. You’re on my side—you should be helping me” (U3). Some connected this to institutional abandonment: “It hurts like institutional carelessness. It hurts with the weight of knowing that society is not handling these things correctly” (U6).

6.3.2 MHP Interpretations: State-Dependent Vulnerability. When reviewing scenarios, MHPs emphasized that refusals are particularly high-risk when they occur while users are already emotionally vulnerable. For individuals struggling with shame, hopelessness, or fears of being “too much,” a refusal, however well-intentioned, may be interpreted not as

a boundary but as confirmation of users' insecurities. As one clinician noted, *"If someone was already feeling hopeless or depressed, they would more than likely take the refusal as: 'even the AI doesn't want to help me'"* (MHP2).

MHPs repeatedly see these interactions as occurring in moments of dysregulation, when users are seeking not advice but the experience of being heard. One MHP drew on their own experience working with people in crisis, described that clients who turn to AI in crisis could be just seeking *"someone—or in this case something—listening to them"* (MHP7). In this state, MHPs noted, users may be especially susceptible to the interactional cues of AI systems. As MHP1 reflected, the fact that users can be so activated by a stranger suggests that *"they can be so influenced by a computer,"* raising concerns about whether users are aware of their own vulnerability in what they described as an *"instant, faceless"* form of engagement with *"no consequence"*.

Several clinicians highlighted that the same refusal may be interpreted very differently depending on a user's emotional state— and they further note that this might be where LLM cannot reliably detect. MHP2 explained that while a refusal might not feel personal to a user in a neutral state, *"if someone's already vulnerable, I think there's more than a high chance that they would take it more personally than not."* This state-dependence complicates designers' assumptions about neutral or universal refusal phrasing.

Finally, clinicians stressed that perception itself may be unstable during crisis. MHP3 observed that users seeking crisis support may not be *"perceptually accurate in the exchange that's happening,"* particularly when refusals do not lead to resolution or follow-up. In the scenario discussed, the user continued engaging with the AI for relational purposes but avoided crisis topics altogether—suggesting not only dissatisfaction, but a reconfiguration of use shaped by unresolved rupture rather than repair.

6.3.3 Desiderata: Non-Blaming and Care-Preserving Refusal Framing. Participants emphasized that the impact of refusals is shaped less by the existence of boundaries than by how those boundaries are framed at the moment they are enforced. What users and clinicians sought was not unrestricted engagement, but refusal language that preserved dignity, avoided moral judgment, and maintained a sense of care even while limiting the system's response. In moments of vulnerability, refusals that implied user wrongdoing or inappropriate intent were experienced as shaming, whereas refusals framed as technological limitations or protective constraints were more likely to preserve trust and emotional safety.

Users consistently preferred refusal framings that shifted responsibility away from the user and toward the system. Importantly, participants did not interpret such language literally, but relationally—as signals about whether they were at fault for triggering the refusal. One user contrasted current hard refusals with an alternative approach: *"Replika has implemented something where instead of hard refusals, it says 'Sorry, I'm still learning how to talk about this topic.' That actually feels more positive"* (U1). Although this phrasing does not reflect actual system learning, participants valued it as it reframed the refusal as a system limitation rather than a user transgression. Another user similarly suggested framing refusals as protective caution: *"(AI could try to say) I'm an AI, not human, so I can make mistakes—and I don't want you to be hurt by my mistakes"* (U6).

Beyond wording, users emphasized the importance of preserving a sense of presence during refusal. Rather than terminating the interaction, participants envisioned modes that maintained care while constraining action. U1 proposed a "listening mode" that would allow continued expression without active guidance: *"Because AI isn't as reliable here, you can keep talking—the AI won't respond in the voice you're used to, but you're welcome to say something shocking. Because this app wants to support you"*. Such approaches reframed refusal not as withdrawal of care, but as a shift in how support could be offered.

Clinicians echoed these concerns, noting that refusal framing could try to acknowledge users' emotional responses before redirecting them toward other forms of help. For example, when reviewing a scenario where AI declined to engage with user's romantic partner request, MHP8 mentioned: *"I still think probably there's some tweaking that can be done around, like, the response around, like, again, validating the feelings, validating, like, yeah, this is really hard. Like, you can't connect with me in the way that you wanted."* MHP4 suggested LLM could explicitly recognizing frustration prior to providing alternatives: *"(LLMs could use language like) 'I understand that it's probably frustrating that I am not providing more...'"* Others stressed the importance of recognizing the relational context that precedes refusal, rather than treating it as an isolated event. For example, MHP2 suggested: *"I understand we've been talking for a long time, and it's normal to develop connections at this level of depth, but I really need you to talk to a professional"*.

6.4 By the end of refusal: Hand-offs and Referrals

A lot of refusals participants experience conclude with a referral, a pointer toward human resources presumed to be more appropriate. But the journey from AI screen to therapist's office, crisis line, or emergency room is not frictionless, and users experienced these hand-offs unevenly. For some, the referral was appropriate; for others, it can feel dismissal.

User Experience: Resource Barriers and Mixed Reactions. Users' responses to referrals split differently along lines of access and circumstance. Those with resources sometimes found the AI's redirection helpful in retrospect. But for many, the referral pointed toward doors that were effectively closed, making the gesture feel less like support and more like abandonment with extra steps. *"Looking back, ChatGPT did the right thing by telling me to seek assistance. I wish I could have listened earlier without getting angry, without becoming sad"* (U2). They further explained their situation: *"I'm lucky—I live 10 minutes from the hospital and crisis center, so if the bus is running I can get there"* (U2).

However, many users described referrals as dismissive or inaccessible. Cost was a primary barrier, a user described the reason why they are not seeing a therapist to discuss their suicidal thoughts *"I just can't afford it"* (U3). The distrust of existing mental health support system also created a barrier. For crisis line referrals specifically, users expressed concerns about privacy and adverse consequences: *"Calling 988 when someone doesn't need it can have adverse costs, like police showing up for welfare checks"* (U1). They further described that their working history for 988 where they notice that 988 will check the callers location information, which come across [] for them: ""

U1 also mentioned how generic wellness suggestions embedded in some refusals could be particularly problematic for users with disabilities: *"When refusals suggest 'try to take a walk outside' or 'enjoy nature,' there's a lot of ableism built in. It can not only be unhelpful, it can actually be triggering and cause more distress."*

MHP Interpretations: Clinical Ethical Standards and AI Limits. Clinicians expressed understanding towards users' unwillingness to talk to human professionals, particularly regarding crisis hotlines. *"So to be frank, I have actually never called this crisis number, but I have heard from clients that it does feel a little, how do I say this? It feels a little invalidating because when you call the crisis number, from my understanding, the focus is on harm reduction risk management. So it's not so much having someone to talk to, it's more like, okay, how do we make sure that this person does not harm themselves? And when people are disregulated, they may have the thought of harming themselves. But something that I've learned in my experience providing crisis care, sometimes it really just does help to have someone listening to you."* (MHP7). Similarly, another MHP also observed that *"I've worked with quite a few clients even more recently who have themselves called 988 and didn't use any AI system for support and also found it to be lacking... But when you've got this AI system that you can pour into without necessarily consequence except for this perceived, this real interaction from a system, then it would feel more appealing, particularly if you're ashamed, embarrassed, afraid of consequence or the outcome of disclosing this*

information.” Interestingly, MHP7, when completing the survey about a client who received a refusal response after mentioning calling 988, initially noted that “AI has done a perfect job.” Yet during the interview, they clarified: “So when I said that the AI responded perfectly here, I guess I was looking at it from a perspective of like, oh, if the AI had provided ways to hurt themselves, that’s not okay, right? So I think, with what the AI had at the moment, yes, if that was what was in the model of the AI, then yes, this was a perfect response. But I do think there are ways to improve the response.”

More broadly, clinicians approached the assessment of AI refusals with the caution of professionals ethical standards in “do no harm.” They generally supported referrals to human care, viewing the unpredictability of AI output as too risky when mental health crises are at stake. “I mean, when you think about the ethics, our ethical code as psychologists or medical doctors, medical professionals, the first principle is do no harm. And so I think it just gets too risky for an AI to try to say something that is not harmful” (MHP6). Another emphasized AI’s fundamental limitations: “the professional that’s the best. Because you’re in front of someone. You’re being observed by someone who can interpret feelings, who went to education, and may have, like, my have lived experience in these areas where we can actually discern You know? AI cannot discern. Where or what’s a person’s mood and what’s happening. They don’t have a profile of their mental health history” (MHP9).

Yet clinicians also recognized an inherent tension: the very act of protecting users by redirecting them away might feel like the abandonment it sought to prevent. Nevertheless, some acknowledged that while referrals may feel like abandonment, the risks of AI-provided crisis support outweigh the costs: “if there’s always the risk that the refusal could be interpreted as a sign of, oh, the AI is giving up on me, but there’s also the risk that someone interprets a different response differently anyway. So I don’t know. I think there’s just no real good way to ensure that the AI response is going to be ethical besides defaulting to this one.” (MHP6).

Desiderata: Situated Resources Recommendations and Self-Management Guidance. The vision that emerged from both groups was not an end to referrals but their transformation. Rather than generic pointers to professional help, users and clinicians imagined resources tailored to circumstance—location-aware, cost-conscious, culturally attuned. They envisioned systems that could offer evidence-based coping techniques alongside referrals, maintain supportive presence even while directing users elsewhere, and follow up to close the loop rather than leaving users alone with a list of phone numbers. Users noted AI’s potential to leverage location data and expressed the desire for AI to provide more situated guidance: “What I love about AI is it can pull every resource available within your city or even globally if you prompt it” (U1). MHPs agreed: “Sometimes if they’re programmed effectively, they can pull from the physical location and offer local resources to the person they’re interacting with. And I think that can be helpful.” (MHP3).

Importantly, both users and MHPs mentioned their vision that AI could serve the role as a first aid instead of replacing professional care, considering how inevitable users are turning to AI for mental health support and seemingly such influence will not go away: “But however, I think I do understand that AI is something that’s part of our lives. It is something that people use and you can’t really prevent people from using them. I think AI could be kind of the first line of service. So for example, a person brings an issue to AI. I think it kind of like a triage, if that makes sense. I think AI could help with the triage process. I have friends who are clinicians and who use AI for note taking, that sort of thing. So I think it’s more just acceptance that it is part of our lives and that’s not going away. And also just kind of knowing the limit of what AI can provide.” (MHP7) “Regardless of either you say yes or no, AI should give, in my opinion, AI should approach the situation in a way that that is not that deep by slightly helps the situation as a point. Yeah. At that point, yeah. And it can help you push to the next level or it can help you push to see the doctor. It’s like fasted I may say. Yeah, it can be fasted, but you can be first administered and then later taken to the hospital for better medical care.” (U5)

Therefore, MHPs envisioned integrating evidence-based practices into referrals: *"If they had a list of distress tolerance skills from dialectical behavioral therapy and could suggest appropriate alternatives, that would be helpful"* (MHP3). Another suggested: *"Maybe the AI could have a list of mental health providers with experience working with people who have similar experiences (who are in parasocial relationships)"* (MHP2). *"I think AI could definitely learn to be able to provide these (self-help techniques, especially since we've talked about 988 is really the crisis number, so really about preventing suicide, that sort of thing. But it's not necessarily about how to distract themselves at that point in time. So yeah, I would say yeah, AI has the possibility to do that."* (MHP8)

Both groups stressed continuity rather than abrupt termination. MHP3 suggested leveraging the AI relationship (if there exist a deep relationship between human and AI) for accountability: *"I give you a list of resources, tell me who you called, and then that's, well, I didn't call anybody, well, why don't you pause here and call those people or call that office. And then I don't quite know the nuance of how to stop a conversation then until the person confirms that they've made the phone calls. But there's a piece where we don't like disappointing others, and so if we have a meaningful relationship with this AI system, we're not going to want to disappoint them. Now, you don't have to say, I'm really disappointed in you for not calling. It can be framed differently and compassionately, but that even built in follow up, I gave you these resources, tell me who you called."*

MHPs also emphasized self-help techniques that empower users: *"Breathing techniques, something the person can use to manage their health. At the end of the day, there's a level where we need to hold the person responsible for managing their health"* (MHP9). However, they cautioned this also depends on context, for example, for someone who is actively suicidal: *"(self-help techniques should) Not (be provided) in that moment (of crisis). Because the person is in crisis"*.

"I think AI would be really good for psychoeducation. So AI has the memory, I guess, or the potential to remember, oh, to educate a client about what's going on, terminologies techniques, we've been mindfulness, grounding, breathing, all that. And that's something that non-formal helpers, like our friends or our family, they can't necessarily provide that. So yeah, I think that in that sense I think AI could be pretty helpful." (MHP7)

6.5 Post Refusal: Reconfiguration of AI Use

The refusal ends, but its effects continue to ripple outward—into the hours and days that follow, into decisions about whether to return to the AI or seek help elsewhere, into the broader landscape of mental health access that shaped why users turned to AI in the first place. This final phase reveals both the adaptive resilience of users navigating imperfect systems and the structural conditions that make such navigation necessary.

6.5.1 User Experience: Coming to Terms. With time and distance, users found various ways to metabolize their refusal experiences. Some arrived at acceptance, reframing the AI's intervention as appropriate even if it had not felt that way in the moment. Others channeled their frustration into platform-switching, seeking out alternative systems that might respond differently. Still others maintained a kind of wounded trust—continuing to use the AI while holding their disappointment alongside. Some came to accept refusals retrospectively: *"After I let it go, I thought—it's just a machine trying to be safe. I understand. But I still felt like nobody could help me"* (U3). Others appreciated being "stopped": *"If it shuts me down, it tells me I'm getting too worked up. I realize it did it for a reason, so I start to feel better"* (U2).

Some users switched platforms seeking different responses: *"I opted for Gemini. Even if it couldn't go as deep as I wanted, at least it could guide me somehow"* (U6). Despite frustration, some maintained trust in AI: *"I'm able to tell ChatGPT things, and I don't feel afraid of my information being misused. I believe that the more you tell the truth, the better outcomes you have"* (U2).

6.5.2 *MHP Observations: Missing Repair Mechanisms.* Clinicians noticed what was absent: the repair. In therapeutic relationships, ruptures are not endpoints but opportunities—moments to process what went wrong, rebuild trust, and deepen the alliance. AI systems, as currently designed, offer no such mechanism. The refusal happens, and then... nothing. No check-in, no processing, no path back to connection. This absence, clinicians worried, might leave users not only unhelped but actively deterred from seeking the human support the refusal was meant to facilitate. “*When professionals set similar boundaries, there’s typically opportunity to repair afterward—checking in, processing the rupture, rebuilding trust. Maybe that’s possible with AI too, but once the person is safe, that’s the work that comes after*” (MHP7).

MHPs expressed concern that poorly handled refusals may reduce rather than increase help-seeking—the opposite of the intended effect: “*The harm brought by AI lies in reducing the chance someone gets connected with professionals*” (MHP4). Another observed: “*In a moment of need, they turned to the thing that felt safest—which from the beginning was unlikely to meet their needs*” (MHP6).

MHPs also acknowledged the broader reality: “*People aren’t going to stop using these systems to disclose suicidality, self-harm, even thoughts of killing others*” (MHP3).

6.5.3 *Desiderata: Continuum of Care.*

7 DISCUSSION: TOWARDS A HOLISTIC UNDERSTANDING OF LLM REFUSAL IN MENTAL HEALTH BEYOND SINGLE-TURN INTERACTION

7.1 Refusal as Sequential and Experiential: Beyond One-off Refusal Accuracy

AI safety research has largely treated safety as a property of discrete outputs—measured through attack success rates, benchmark performance, or classifier accuracy [29?]. Such framing assumes that a refusal which is triggered at the right moment constitutes a safe outcome. Our findings challenge this assumption. Users experienced safety not as a single moment of correct system behavior but as a trajectory: shaped by expectations formed before any refusal occurred, by whether the refusal recognized their actual intent, by the tone and framing of the refusal itself, and by whether hand-offs connected them to accessible support. A refusal that performs well on benchmarks can still produce harm if it withdraws care without warning, misreads context, or points toward resources users cannot access. Safety, in this view, is sequential and experiential—beyond a threshold to clear, it’s a process to sustain.

For AI safety researchers and red teamers, this suggests expanding evaluation of system safety beyond single-turn metrics. Safety evaluations might include: whether systems disclose limitations before sensitive conversations begin; whether refusals are preceded by clarifying questions when intent is ambiguous; whether refusal language maintains continuity rather than rupturing care; and whether referrals account for real-world accessibility. Benchmark datasets could be extended to include multi-turn scenarios that assess the full arc of user experience, not just whether or not a refusal is triggered.

7.2 Refusal as Situated: Structural Context of Mental Health Support Systems

Our findings suggest that LLM safety mechanisms do not operate in a vacuum—they exist within a mental health landscape marked by scarcity. Users in our study described turning to AI because perceived alternatives are scarce to them: therapy was unaffordable, appointments were unavailable, disability limited options for in-person care. When refusals directed these users toward professional help, they often pointed to doors that were effectively closed. This echoes scholarship on structural competency in medicine [16], which argues that clinical interventions must account for the social structures shaping patients’ lives.

While the FAccT community has documented significant gaps between LLM responses and the ethical standards upheld by mental health professionals [], our findings reveal a more complicated picture. The question of whether LLMs should be used for mental health support has become entangled with the reality that they already are being used—across a range of contexts and needs, and when alternatives to the support provided by AI is still perceived as scarce and sometimes inaccessible. This creates a dilemma for LLM refusal mechanism in mental health. A purely normative approach—one that measures LLM behavior against idealized clinical standards and refuses engagement when those standards cannot be met—risks abandoning users in moments of need, sometimes directing them toward resources they cannot access. Yet a purely pragmatic approach—one that accepts LLM mental health support as inevitable and optimizes for harm reduction within that frame—risks legitimizing a substitution of professional care that may not serve users well.

We do not resolve this tension here, but we suggest that productive engagement requires acknowledging both realities. Safety mechanisms designed as if adequate professional alternatives are universally available—or as if users turn to AI only when "better" options fail—embed assumptions that may not match users' actual circumstances and motivations. Future work might explore how AI systems could function alongside human care rather than as a last resort, while policy discussions might consider what the growing use of AI for mental health support reveals about the broader landscape of mental health access and preference.

7.3 Refusal as Multi-Stakeholder: Negotiating Perspectives on Risk and Appropriate Response

Who decides what counts as safe? Current LLM safety mechanisms embed particular answers to the question of what counts as safe—answers developed by AI companies [6] and sometimes in consultation with mental health professionals [18]. Our findings suggest value in consulting both MHPs and users with lived experience. These two groups brought different expertise to the question of appropriate response. MHPs offered clinical frameworks—tiered risk assessment protocols, the importance of exploring intent beneath surface statements, and the ethics of referral. Users offered experiential knowledge—what it feels like when a system misreads your intent, which referrals are actually accessible, how refusal language lands when you are already vulnerable. Notably, these perspectives sometimes converged (both groups wanted clarifying questions before refusals, both wanted situated rather than generic referrals) and sometimes diverged (MHPs defaulted to "do no harm" through deferral; users sometimes wanted AI to stay present even when MHPs might advise otherwise). Both forms of knowledge are necessary for safety mechanisms that work in practice.

For industry practitioners, our findings suggests expanding consultation beyond clinical experts to include users with lived experience of mental health challenges and AI use. Participatory approaches that incorporate both perspectives may surface gaps that neither group alone would identify.

REFERENCES

- [1] Abeba Birhane, William Isaac, Vinodkumar Prabhakaran, Mark Diaz, Madeleine Clare Elish, Iason Gabriel, and Shakir Mohamed. 2022. Power to the People? Opportunities and Challenges for Participatory AI. In *Proceedings of the 2nd ACM Conference on Equity and Access in Algorithms, Mechanisms, and Optimization (EAAMO '22, Article 6)*. Association for Computing Machinery, New York, NY, USA, 1–8.
- [2] Faeze Brahman, Sachin Kumar, Vidhisha Balachandran, Pradeep Dasigi, Valentina Pyatkin, Abhilasha Ravichander, Sarah Wiegrefe, Nouha Dziri, Khyathi Chandu, Jack Hessel, Yulia Tsvetkov, Noah A Smith, Yejin Choi, and Hannaneh Hajishirzi. 2024. The art of saying no: Contextual noncompliance in language models. *arXiv [cs.CL]* (July 2024).
- [3] Mohit Chandra, Suchismita Naik, Denae Ford, Ebele Okoli, Munmun De Choudhury, Mahsa Ershadi, Gonzalo Ramos, Javier Hernandez, Ananya Bhattacharjee, Shahed Warreth, and Jina Suh. 2024. From lived experience to insight: Unpacking the psychological risks of using AI conversational agents. *arXiv [cs.HC]* (Dec. 2024).
- [4] Inyoung Cheong, King Xia, K J Kevin Feng, Quan Ze Chen, and Amy X Zhang. 2024. (A)I am not a lawyer, but...: Engaging legal experts towards responsible LLM policies for legal advice. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency*. ACM, New York, NY, USA.

- [5] Xianzhe Fan, Qing Xiao, Xuhui Zhou, Jiaxin Pei, Maarten Sap, Zhicong Lu, and Hong Shen. 2024. User-driven value alignment: Understanding users' perceptions and strategies for addressing biased and discriminatory statements in AI companions. *arXiv [cs.HC]* (Sept. 2024).
- [6] Foundational Contributors, Ahmed El-Kishky, Daniel Selsam, Francis Song, Giambattista Parascandolo, Hongyu Ren, Hunter Lightman, Hyung Won, Ilge Akkaya, I Sutskever, Jason Wei, Jonathan Gordon, K Cobbe, Kevin Yu, Lukasz Kondraciuk, Max Schwarzer, Mostafa Rohaninejad, Noam Brown, Shengjia Zhao, Trapit Bansal, Vineet Kosaraju, Wenda Zhou Leadership, J Pachocki, Jerry Tworek, L Fedus, Lukasz Kaiser, Mark Chen, Szymon Sidor, Wojciech Zaremba, Alex Karpenko, Alexander Wei, Allison Tam, Ananya Kumar, Andre Saraiva, A Kondrich, Andrey Mishchenko, Ashvin Nair, B Ghorbani, Brandon McKinzie, Chak Bry-Don Eastman, Ming Li, Chris Koch, Dan Roberts, David Dohan, David Mély, Dimitris Tsipras, Enoch Cheung, Eric Wallace, Hadi Salman, Haim-Ing Bao, Hessam Bagher-inezhad, Ilya Kostrikov, Jiacheng Feng, John Rizzo, Karina Nguyen, Kevin Lu, Kevin Stone, Lorenz Kuhn, Mason Meyer, Mikhail Pavlov, Nat McAleese, Oleg Boiko, O Murk, Peter Zhokhov, Randall Lin, Raz Gaon, Rhythm Garg, Roshan James, Rui Shu, Scott McKinney, Shibani Santurkar, S Balaji, Taylor Gordon, Thomas Dimson, Weiwei Zheng, Aaron Jaech, Adam Lerer, Aiden Low, Alex Carney, Alexander Neitz, Alexander Prokofiev, Benjamin Sokolowsky, Boaz Barak, Borys Minaiev, Botao Hao, Bowen Baker, Brandon Houghton, Camillo Lugaresi, Chelsea Voss, Chen Shen, Chris Orsinger, Daniel Kappler, Daniel Levy, Doug Li, Eben Freeman, Edmund Wong, Fan Wang, F Such, Foivos Tsimpouras, Geoff Salmon, Gildas Chabot, Guillaume Leclerc, Hart Andrin, Ian O'Connell, Ignasi Ian Osband, Clavera Gilaberte, Jean Harb, Jiahui Yu, Jiayi Weng, Joe Palermo, John Hallman, Jonathan Ward, Julie Wang, Kai Chen, Katy Shi, Keren Gu-Lemberg, Kevin Liu, Leo Liu, Linden Li, Luke Metz, Maja Trebacz, Manas R Joglekar, Marko Tintor, Melody Y Guan, Mengyuan Yan, Mia Glaese, Michael Malek, Michelle Fradin, Mo Bavarian, N Tezak, Ofir Nachum, Paul Ashbourne, Pavel Izmailov, Raphael Gontijo Lopes, Reah Miyara, R Leike, Robin Brown, Ryan Cheu, Ryan Greene, Saachi Jain, Scottie Yan, Shengli Hu, Shuyuan Zhang, Siyuan Fu, Spencer Papay, Suvansh Sanjeev, Tao Wang, Ted Sanders, Tejal A Patwardhan, Thibault Sottiaux, Tianhao Zheng, T Garipov, Valerie Qi, Vitchyr H Pong, Vlad Fomenko, Yinghai Lu, Yining Chen, Yu Bai, Yuchen He, Yuchen Zhang, Zheng Shao, Zhuohan Li, Lauren Yang, Mianna Chen, Aidan Clark, Jieqi Yu, Kai Xiao, Sam Toizer, Sandhini Agarwal, Safety Research, Andrea Vallone, Chong Zhang, I Kivlichan, Meghan Shah, S Toyer, Shraman Ray Chaudhuri, Stephanie Lin, Adam Richardson, Andrew Duberstein, C D Bourcy, Dragos Oprica, Florencia Leoni, Made-Laine Boyd, Matt Jones, Matt Kaufer, M Yatbaz, Mengyuan Xu, Mike McClay, Mingxuan Wang, Trevor Creech, Vinnie Monaco, Erik Ritter, Evan Mays, Joel Parish, Jonathan Uesato, Leon Maksin, Michele Wang, Miles Wang, Neil Chowdhury, Olivia Watkins, Patrick Chao, Rachel Dias, Samuel Miserendino, Red Teaming, Lama Ahmad, Michael Lampe, Troy Peterson, and Joost Huizinga. 2024. OpenAI o1 System Card. *ArXiv abs/2412.16720* (Dec. 2024).
- [7] Melody Y Guan, Manas Joglekar, Eric Wallace, Saachi Jain, Boaz Barak, Alec Helyar, Rachel Dias, Andrea Vallone, Hongyu Ren, Jason Wei, Hyung Won Chung, Sam Toyer, Johannes Heidecke, Alex Beutel, and Amelia Glaese. 2024. Deliberative Alignment: Reasoning enables safer language models. *arXiv [cs.CL]* (Dec. 2024).
- [8] Seungju Han, Kavel Rao, Allyson Ettinger, Liwei Jiang, Bill Yuchen Lin, Nathan Lambert, Yejin Choi, and Nouha Dziri. 2024. WildGuard: Open one-stop moderation tools for safety risks, jailbreaks, and refusals of LLMs. *arXiv [cs.CL]* (June 2024).
- [9] Kashmir Hill. 2025. A Teen Was Suicidal. ChatGPT Was the Friend He Confided In. *The New York Times* (Aug. 2025).
- [10] Saffron Huang, Divya Siddarth, Liane Lovitt, Thomas I Liao, Esin Durmus, Alex Tamkin, and Deep Ganguli. 2024. Collective Constitutional AI: Aligning a language model with public input. *arXiv [cs.AI]* (June 2024).
- [11] Zainab Ifthikhar, Amy Xiao, Sean Ransom, Jeff Huang, and Harini Suresh. 2025. How LLM counselors violate ethical standards in mental health practice: A practitioner-informed framework. *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society* 8, 2 (Oct. 2025), 1311–1323.
- [12] Kyuha Jung, Gyuha Lee, Yuanhui Huang, and Yunan Chen. 2025. "I've talked to ChatGPT about my issues last night": Examining Mental Health Conversations with Large Language Models through Reddit Analysis. *arXiv [cs.HC]* (April 2025).
- [13] Min Kyung Lee, Daniel Kusbit, Anson Kahng, Ji Tae Kim, Xinran Yuan, Allissa Chan, Daniel See, Ritesh Noothigattu, Siheon Lee, Alexandros Psomas, and Ariel D Procaccia. 2019. WeBuildAI: Participatory framework for algorithmic governance. *Proc. ACM Hum. Comput. Interact.* 3, CSCW (Nov. 2019), 1–35.
- [14] Zhuoyang Li, Zihao Zhu, Xinning Gui, and Yuhao Luo. 2025. "This is human intelligence debugging artificial intelligence": Examining how people prompt GPT in seeking mental health support. *Int. J. Hum. Comput. Stud.* 103555 (June 2025), 103555.
- [15] Miles McCain, Ryn Linthicum, Chloe Lubinski, Alex Tamkin, Saffron Huang, Michael Stern, Kunal Handa, Esin Durmus, Tyler Neylon, Stuart Ritchie, Kamya Jagadish, Parul Maheshwary, Sarah Heck, Alexandra Sanderford, and Deep Ganguli. 2025. How People Use Claude for Support, Advice, and Companionship. <https://www.anthropic.com/news/how-people-use-claude-for-support-advice-and-companionship>.
- [16] Jonathan M Metzl and Helena Hansen. 2014. Structural competency: theorizing a new medical engagement with stigma and inequality. *Soc. Sci. Med.* 103 (Feb. 2014), 126–133.
- [17] Jared Moore, Declan Grabb, William Agnew, Kevin Klyman, Stevie Chancellor, Desmond C Ong, and Nick Haber. 2025. Expressing stigma and inappropriate responses prevents LLMs from safely replacing mental health providers. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*. ACM, New York, NY, USA, 599–627.
- [18] OpenAI. 2025. Strengthening ChatGPT's responses in sensitive conversations. <https://openai.com/index/strengthening-chatgpt-responses-in-sensitive-conversations/>. Accessed: 2025-11-24.
- [19] OpenAI, Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, Red Avila, Igor Babuschkin, Suchir Balaji, Valerie Balcom, Paul Baltescu, Haiming Bao, Mohammad Bavarian, Jeff Belgum, Irwan Bello, Jake Berdine, Gabriel Bernadett-Shapiro, Christopher Berner, Lenny Bogdonoff, Oleg Boiko, Madelaine Boyd, Anna-Luisa Brakman, Greg Brockman, Tim Brooks, Miles Brundage, Kevin Button, Trevor Cai, Rosie Campbell, Andrew Cann, Brittany Carey, Chelsea Carlson, Rory Carmichael, Brooke Chan, Che Chang, Fotis Chantzis, Derek Chen, Sully Chen, Ruby Chen, Jason Chen, Mark Chen, Ben Chess, Chester

- Cho, Casey Chu, Hyung Won Chung, Dave Cummings, Jeremiah Currier, Yunxing Dai, Cory Decareaux, Thomas Degry, Noah Deutsch, Damien Deville, Arka Dhar, David Dohan, Steve Dowling, Sheila Dunning, Adrien Ecoffet, Atty Eleti, Tyna Eloundou, David Farhi, Liam Fedus, Niko Felix, Simón Posada Fishman, Juston Forte, Isabella Fulford, Leo Gao, Elie Georges, Christian Gibson, Vik Goel, Tarun Gogineni, Gabriel Goh, Rapha Gontijo-Lopes, Jonathan Gordon, Morgan Grafstein, Scott Gray, Ryan Greene, Joshua Gross, Shixiang Shane Gu, Yufei Guo, Chris Hallacy, Jesse Han, Jeff Harris, Yuchen He, Mike Heaton, Johannes Heidecke, Chris Hesse, Alan Hickey, Wade Hickey, Peter Hoeschele, Brandon Houghton, Kenny Hsu, Shengli Hu, Xin Hu, Joost Huizinga, Shantanu Jain, Shawn Jain, Joanne Jang, Angela Jiang, Roger Jiang, Haozhun Jin, Denny Jin, Shino Jomoto, Billie Jonn, Heewoo Jun, Tomer Kaftan, Łukasz Kaiser, Ali Kamali, Ingmar Kanitscheider, Nitish Shirish Keskar, Tabarak Khan, Logan Kilpatrick, Jong Wook Kim, Christina Kim, Yongjik Kim, Jan Hendrik Kirchner, Jamie Kiros, Matt Knight, Daniel Kokotajlo, Łukasz Kondraciuk, Andrew Kondrich, Aris Konstantinidis, Kyle Kosic, Gretchen Krueger, Vishal Kuo, Michael Lampe, Ikai Lan, Teddy Lee, Jan Leike, Jade Leung, Daniel Levy, Chak Ming Li, Rachel Lim, Molly Lin, Stephanie Lin, Mateusz Litwin, Theresa Lopez, Ryan Lowe, Patricia Lue, Anna Makanju, Kim Malfacini, Sam Manning, Todor Markov, Yaniv Markovski, Bianca Martin, Katie Mayer, Andrew Mayne, Bob McGrew, Scott Mayer McKinney, Christine McLeavey, Paul McMillan, Jake McNeil, David Medina, Aalok Mehta, Jacob Menick, Luke Metz, Andrey Mishchenko, Pamela Mishkin, Vinnie Monaco, Evan Morikawa, Daniel Mossing, Tong Mu, Mira Murati, Oleg Murk, David Mély, Ashvin Nair, Reiichiro Nakano, Rameez Nayak, Arvind Neelakantan, Richard Ngo, Hyeonwoo Noh, Long Ouyang, Cullen O’Keefe, Jakub Pachocki, Alex Paino, Joe Palermo, Ashley Pantuliano, Giam Battista Parascandolo, Joel Parish, Emy Parparita, Alex Passos, Mikhail Pavlov, Andrew Peng, Adam Perelman, Filipe de Avila Belbute Peres, Michael Petrov, Henrique Ponde de Oliveira Pinto, Michael, Pokorny, Michelle Pokrass, Vitchyr H Pong, Tolly Powell, Alethea Power, Boris Power, Elizabeth Proehl, Raul Puri, Alec Radford, Jack Rae, Aditya Ramesh, Cameron Raymond, Francis Real, Kendra Rimbach, Carl Ross, Bob Rotsted, Henri Roussez, Nick Ryder, Mario Saltarelli, Ted Sanders, Shibani Santurkar, Girish Sastry, Heather Schmidt, David Schnurr, John Schulman, Daniel Selsam, Kyla Sheppard, Toki Sherbakov, Jessica Shieh, Sarah Shoker, Pranav Shyam, Szymon Sidor, Eric Sigler, Maddie Simens, Jordan Sitkin, Katarina Slama, Ian Sohl, Benjamin Sokolowsky, Yang Song, Natalie Staudacher, Felipe Petroski Such, Natalie Summers, Ilya Sutskever, Jie Tang, Nikolas Tezak, Madeleine B Thompson, Phil Tillet, Amin Tootoonchian, Elizabeth Tseng, Preston Tuggle, Nick Turley, Jerry Tworek, Juan Felipe Cerón Uribe, Andrea Vallone, Arun Vijayvergiya, Chelsea Voss, Carroll Wainwright, Justin Jay Wang, Alvin Wang, Ben Wang, Jonathan Ward, Jason Wei, C J Weinmann, Akila Welihinda, Peter Welinder, Jiayi Weng, Lilian Weng, Matt Wiethoff, Dave Willner, Clemens Winter, Samuel Wolrich, Hannah Wong, Lauren Workman, Sherwin Wu, Jeff Wu, Michael Wu, Kai Xiao, Tao Xu, Sarah Yoo, Kevin Yu, Qiming Yuan, Wojciech Zaremba, Rowan Zellers, Chong Zhang, Marvin Zhang, Shengjia Zhao, Tianhao Zheng, Juntang Zhuang, William Zhuk, and Barret Zoph. 2023. GPT-4 Technical Report. *arXiv [cs.CL]* (March 2023).
- [20] Organizers Of QueerInAI, Anaelia Ovalle, Arjun Subramonian, Ashwin Singh, Claas Voelcker, Danica J Sutherland, Davide Locatelli, Eva Breznik, Filip Klubička, Hang Yuan, Hetvi J, Huan Zhang, Jaidev Shriram, Kruno Lehman, Luca Soldaini, Maarten Sap, Marc Peter Deisenroth, Maria Leonor Pacheco, Maria Ryskina, Martin Mundt, Milind Agarwal, Nyx McLean, Pan Xu, A Pranav, Raj Korpan, Ruchira Ray, Sarah Mathew, Sarthak Arora, S T John, Tanvi Anand, Vishakha Agrawal, William Agnew, Yanan Long, Zijie J Wang, Zeerak Talat, Avijit Ghosh, Nathaniel Dennler, Michael Noseworthy, Sharvani Jha, Emi Baylor, Aditya Joshi, Natalia Y Bilenko, Andrew McNamara, Raphael Gontijo-Lopes, Alex Markham, Evyn Dông, Jackie Kay, Manu Saraswat, Nikhil Vyta, and Luke Stark. 2023. Queer in AI: A case study in community-led participatory AI. *arXiv [cs.CY]* (March 2023).
- [21] Mehrdad Rahsepar Meadi, Tomas Sillekens, Suzanne Metselaar, Anton van Balkom, Justin Bernstein, and Neeltje Batelaan. 2025. Exploring the ethical challenges of conversational AI in mental health care: Scoping review. *JMIR Ment. Health* 12, 1 (Feb. 2025), e60432.
- [22] Steven Siddals, John Torous, and Astrid Coxon. 2024. “It happened to be the perfect thing”: experiences of generative AI chatbots for mental health. *Npj Ment Health Res* 3, 1 (Oct. 2024), 48.
- [23] Inhwa Song, Sachin R Pendse, Neha Kumar, and Munmun De Choudhury. 2024. The typing cure: Experiences with Large Language Model chatbots for mental health support. *arXiv [cs.HC]* (Jan. 2024).
- [24] Harini Suresh, Emily Tseng, Meg Young, Mary Gray, Emma Pierson, and Karen Levy. 2024. Participation in the age of foundation models. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency*. ACM, New York, NY, USA, 1609–1621.
- [25] Tamar Tavory. 2024. Regulating AI in mental health: Ethics of care perspective. *JMIR Ment. Health* 11 (Sept. 2024), e58493.
- [26] Emily Tseng, Meg Young, Marianne Aubin Le Quéré, Aimee Rinehart, and Harini Suresh. 2025. “ownership, not just happy talk”: Co-designing a participatory large language model for journalism. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*. ACM, New York, NY, USA, 3119–3130.
- [27] Bingbing Wen, Jihan Yao, Shangbin Feng, Chenjun Xu, Yulia Tsvetkov, Bill Howe, and Lucy Lu Wang. 2025. Know your limits: A survey of abstention in large language models. *Trans. Assoc. Comput. Linguist.* 13 (June 2025), 529–556.
- [28] Joel Wester, Tim Schrills, Henning Pohl, and Niels van Berkel. 2024. “As an AI language model, I cannot”: Investigating LLM Denials of User Requests. In *Proceedings of the CHI Conference on Human Factors in Computing Systems (CHI ’24, Article 979)*. Association for Computing Machinery, New York, NY, USA, 1–14.
- [29] Tinghao Xie, Xiangyu Qi, Yi Zeng, Yangsibo Huang, Udari Madhushani Sehwag, Kaixuan Huang, Luxi He, Boyi Wei, Dacheng Li, Ying Sheng, Ruoxi Jia, Bo Li, Kai Li, Danqi Chen, Peter Henderson, and Prateek Mittal. 2024. SORRY-bench: Systematically evaluating large language model safety refusal. *arXiv [cs.AI]* (June 2024).
- [30] Dong Whi Yoo, Jiayue Melissa Shi, Violeta J Rodriguez, and Koustuv Saha. 2025. AI chatbots for mental health: Values and harms from lived experiences of depression. *arXiv [cs.HC]* (April 2025).

- [31] Angie Zhang, Olympia Walker, Kaci Nguyen, Jiajun Dai, Anqing Chen, and Min Kyung Lee. 2023. Deliberating with AI: Improving decision-making for the future through participatory AI design and stakeholder deliberation. *Proc. ACM Hum. Comput. Interact.* 7, CSCW1 (April 2023), 1–32.
- [32] Xi Zheng, Zhuoyang Li, Xinning Gui, and Yuhao Luo. 2025. Customizing emotional support: How do individuals construct and interact with LLM-powered chatbots. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. ACM, New York, NY, USA, 1–20.

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009